




KAWAN


*It doesn't believe you. **Yet.***

A voiced Live2D accountability companion that verifies your commitments with real evidence — it never does the work for you.

Chutes Hack Malaysia 2026 • Corporate + SIWC tracks



The accountability gap



Every AI coach trusts what you tell it. **None of them check.**

- **Self-report is gameable**

"I finished it" — and no one ever verifies.

- **Slacking is free**

No cost in the moment, so it always wins.

- **One slip and you quit**

A broken streak plus guilt — app abandoned.

Meet Kawan

A loveable-but-skeptical friend who helps you set **one real commitment** — then **checks, with fetched evidence**, that you actually did it.

TWO THINGS IT REFUSES



Self-report

It won't take your word — it fetches the evidence itself.



Doing your work


No answers, ever. Only process, scope, sequence and time.




Win seeds the next commitment; a **miss** routes to cheap recovery and a re-commit — no streak to break.

See it work


The evidence loop, end to end —
real fetches, not self-report.




01 · COMPOSE One sentence, ≤ 60s — then pick a persona.



02 · THE STAGE A living avatar that reacts — skeptical, then pleased.




03 · CHECK NOW It fetches a real GitHub commit and names it.



04 · VERDICT Vision judges it, attests it — then you win.

The agent framework


Five role-specialized agents — not one chatbot.
Structured JSON is the control plane.



The guardrail: the AI writes exactly one table — soft context. Hard fields — deadline, deliverable, status, stakes — are user/system-only and audited. **Kawan can never move the goalposts.**

System architecture

Boring where it should be — **one process, one event loop**. Deep only where it counts.



Frontend → **Vercel** · Backend → **Render** (1 instance) · DB → **Supabase**. No external queue — nothing extra to break on stage.

Built on Chutes

Five judge-visible touchpoints – everything runs on the **user's own confidential compute**.



01

TEE + attestation

Intel TDX confidential inference, surfaced in the UI. Zero prompt logging.



02

Structured outputs

JSON is the control plane – it drives product logic, not just formatting.



03

Multimodal judging

A TEE vision model reads your screenshot and reasons about it.



04

Inline failover

Per-persona primary, failover routing – in a single call.



05

Sign-in-with-Chutes

The user's token is the Bearer. Every call bills **their own balance**.

MODEL POOL

gemma-4-31B-turbo-TEE

DeepSeek-V3.2-TEE

Kimi-K2.6-TEE

Privacy, billing and product logic – one story.

Business value

Self-report · “trust me”



Verified evidence · attested

Impact

- A verifiable commitment in **under 60 seconds**.
- Accountability with **zero gameable self-report**.
- **No streaks** — cutting the abandonment that kills habit apps.

Market

- The AI-accountability wave runs on self-report — **verification is the open gap**.
- Students, knowledge-workers, indie builders.
- **Corporates** for team accountability — the Corporate track.

Viability

- **SIWC = user-funded compute** — operator inference cost $\approx$ \$0 at scale.
- **TEE = enterprise privacy** — operators can't read your goals or train on them.
- Honest, shippable scope.

Everyone else trusts you. Kawan checks — privately, on your own compute.

Roadmap & team

Honest scope is a feature — **we plan and verify**; we don't fake retention in 20 days.

● **SHIPPED NOW** — live in the product

- ✓ SIWC auth
- ✓ 3 personas + Live2D
- ✓ VN / Messages workspace
- ✓ Commitments + timeline
- ✓ Evidence: GitHub · screenshot · file
- ✓ Notifications: push · Telegram · email
- ✓ Achievements + share cards

◉ **NEXT** — clearly labeled vision

- + Monetary stakes
- + Multi-commitment
- + Friend graph / leaderboard
- + Deeper retention

TEAM

4 people — full-stack web & AI / LLM



ZJ
/AlaskanTuna



JH
/c3638



kymil4
/kymil4




JRM
/WhiteAvocad0



KAWAN

*No one keeps a promise to a cheerleader.
They keep it to a friend who **checks**.*

ALL ON CHUTES • TEE inference • Structured outputs • Multimodal judging • User-funded billing



Scan to open Kawan

Repo → [github]

Demo → [video]

Contact → [email]